



Experimental Workflow: RL-Learned Gap Penalties for MSA

Training Loop (Contextual Bandit)



Evaluation



Key insight: The agent learns op ~ 3.0 (2x default) and ep ~ 0.4 (vs default 0.0), yielding +2.0% reward and +5.4% TC improvement